

# Single-Source Shortest Path

Analysis of Algorithms

# Shortest Path Applications

- Map routing
- Seam carving
- Robot navigation
- Texture mapping
- Typesetting in TeX
- Urban traffic planning
- Optimal pipelining of VLSI chip
- Telemarketer operator scheduling
- Routing of telecommunications messages
- Network routing protocols (OSPF, BGP, RIP)
- Exploiting arbitrage opportunities in currency exchange
- Optimal truck routing through given traffic congestion pattern



# Single-Source Shortest Path

- Single-source shortest-path algorithms find the series of edges between two vertices that has the smallest total weight
- A minimum spanning tree algorithm won't work for this because it would skip an edge of larger weight and include many edges with smaller weights that could result in a longer path than the single edge

# Single-Source Shortest Path

- Initialize  $\text{distTo}[\text{source}] = 0$
- Initialize  $\text{distTo}[v] = \infty$  for all other vertices,  $v$
- Optimality condition:
  - For each edge  $(u, v)$ ,  $\text{distTo}[v] \leq \text{distTo}[u] + w(u, v)$
- To achieve the optimal condition, repeat until satisfied:
  - Relax an edge
  - Relaxing an edge means getting “closer to optimal” on each iteration

# Edge Relaxation

- “Relaxing” an edge:
  - If an edge  $(u, v)$  with weight  $w$  gives a shorter path from the source to  $v$  through  $u$ , then update the  $\text{distTo}[v]$  and set the parent (predecessor) of  $v$  to  $u$ :

RELAX( $u, v, W$ ):

    If  $\text{distTo}[v] > \text{distTo}[u] + w[u, v]$

$\text{distTo}[v] := \text{distTo}[u] + w[u, v]$

$\text{parent}[v] := u$

# Dijkstra's Algorithm

- Dijkstra's algorithm is similar to the Prim MST algorithm, but instead of just looking at a ***single*** shortest edge from a vertex to a vertex in the fringe, we look at the ***overall*** shortest path from the start vertex to a vertex in the fringe
- Note: In order for Dijkstra's method to work, all weights must be non-negative

# Dijkstra's Algorithm

DIJKSTRA(Graph, source):

- Initialize distance to every vertex to  $\infty$

- Initialize distance to source to 0

- Initialize shortest path set S to empty

- Insert all vertices into priority queue, PQ

while PQ is not empty:

- $u :=$  extract the vertex with the min value in the PQ


- Insert vertex u into set S

- for each vertex v adjacent to u:


  - RELAX(u, v)

  - update the priority of v

# Dijkstra Example




Initial fringe:






# Dijkstra Example

Select edge A-B:




Select edge A-C:




# Dijkstra Example

Select edge B-E:




Select edge A-F:




# Dijkstra Example

Select edge F-D:



Select edge B-G:



# Dijkstra and Prim



- Dijkstra's shortest path algorithm is essentially the same as Prim's minimum spanning tree algorithm
- The main distinction between the two is the rule that is used to choose next vertex for the tree
  - Prim: Choose the closest vertex (smallest weight) ***to any vertex*** in the minimum spanning tree so far
  - Dijkstra's: Choose the closest vertex (smallest weight) ***from the source vertex***
  - Note: DFS and BFS are also in this family of algorithms

# Analysis of Dijkstra's Algorithm

- Algorithm:
  - While the PQ is not empty, return and remove the “best” vertex (the one closest to the source), and update the priorities of all the neighbors of that best vertex
  - The overall runtime depends on implementation:
    - Using a simple array or linked list causes the runtime to be proportional to  $N^2 + M \approx N^2$  (best for dense graph)
    - Using a binary heap causes the total runtime to be proportional to  $N \log N + M \log N \approx M \log N$  (best for sparse graph)

# Negative Weights

- Dijkstra does not work with negative weights
  - Dijkstra selects vertex 3 immediately after 0, but shortest path from 0 to 3 is  $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$
- What about re-weighting the edges?
  - Add a constant to every edge weight to make all edges positive doesn't work either
  - Adding 9 to each edge weight causes Dijkstra to again incorrectly select vertex 3
- Conclusion: We need a different algorithm for negative weights



# Bellman-Ford Algorithm

BELLMAN-FORD(Graph, source):

Initialize distance to every vertex to  $\infty$

Initialize distance to source to 0

for  $i := 1$  to  $N-1$

    for each edge  $(u, v)$

        RELAX( $u, v$ )

for each edge  $(u, v)$

    if  $\text{distTo}[v] > \text{distTo}[u] + w[u, v]$


        return false

return true


# Bellman-Ford Example

Each pass relaxes the edges  
in some arbitrary order:


$(t, x)$ ,  $(t, y)$ ,  $(t, z)$ ,  $(x, t)$ ,  
 $(y, x)$ ,  $(y, z)$ ,  $(z, x)$ ,  $(z, s)$ ,  
 $(s, t)$ ,  $(s, y)$




Start




After Pass 1



After Pass 2



After Pass 3



After Pass 4



# Bellman-Ford Java Code

```
public class BellmanFordSP
{
    private double[] distTo;
    private DirectedEdge[] edgeTo;
    private boolean[] onQ;
    private Queue<Integer> queue;

    public BellmanFordSPT(EdgeWeightedDigraph G, int s)
    {
        distTo = new double[G.V()];
        edgeTo = new DirectedEdge[G.V()];
        onQ = new boolean[G.V()];
        queue = new Queue<Integer>();

        for (int v = 0; v < V; v++)
            distTo[v] = Double.POSITIVE_INFINITY;
        distTo[s] = 0.0;

        queue.enqueue(s);
        while (!queue.isEmpty())
        {
            int v = queue.dequeue();
            onQ[v] = false;
            for (DirectedEdge e : G.adj(v))
                relax(e);
        }
    }
}
```

queue of vertices whose  
distTo[] value changes

```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
        if (!onQ[w])
        {
            queue.enqueue(w);
            onQ[w] = true;
        }
    }
}
```

# Analysis of Bellman-Ford

- Weights can be negative, but the graph cannot have negative-weight cycles!
- Bellman-Ford will detect a negative-weight cycle
  - Run the algorithm one more iteration: if the shortest path returned is ***less than*** the shortest path from the previous iteration, then return false (no solution exists because of a negative-weight cycle)
  - Else return true (the path returned is the shortest path solution)
- Runtime
  - $N-1$  passes, each pass looks at  $M$  edges
  - Thus, the total runtime is proportional to  **$N \cdot M$**

# Analysis of Bellman-Ford

- Bellman-Ford is naturally distributed, whereas Dijkstra is naturally local
- Can be used for a network routing protocol
  - Change from a source-driven algorithm to a destination-driven algorithm by just reversing the direction of the edges in Bellman-Ford
  - Change to a “push-based” algorithm: as soon as a vertex  $v$  discovers it's shortest path to the destination,  $v$  notifies all of its neighbors
    - This works well even in an asynchronous network

# Acyclic Shortest Path Algorithm

- Suppose an edge-weighted digraph has no directed cycles (i.e., it is a weighted DAG)
- Consider the vertices in topological order
- Relax all edges pointing from that vertex

DAG-SHORTEST-PATHS( $G$ , source):

Topologically sort the vertices of  $G$

Initialize distance to every vertex to  $\infty$


Initialize distance to source to 0

for each vertex  $u$  taken in topological order

for each vertex  $v$  adjacent to  $u$


RELAX( $u$ ,  $v$ )

# Acyclic Shortest Path Algorithm



First, topologically sort the vertices (assume source is  $s$ ). This figure shows after the first iteration of the for loop. The colored vertex,  $r$ , was used as  $u$  in this iteration.

# Acyclic Shortest Path Algorithm




After the second iteration of the for loop.

The colored vertex, s, was used as  $u$  in this iteration.

The bold edges indicate the shortest path from source.

# Acyclic Shortest Path Algorithm




After the third iteration of the for loop.

The colored vertex,  $t$ , was used as  $u$  in this iteration.

The bold edges indicate the shortest path from source.

# Acyclic Shortest Path Algorithm




After the fourth iteration of the for loop.

The colored vertex, x, was used as  $u$  in this iteration.

The bold edges indicate the shortest path from source.



# Acyclic Shortest Path Algorithm




After the fifth iteration of the for loop.

The colored vertex, y, was used as  $u$  in this iteration.

The bold edges indicate the shortest path from source.

# Acyclic Shortest Path Algorithm



After the sixth iteration of the for loop (final values).  
The colored vertex,  $z$ , was used as  $u$  in this iteration.  
The bold edges indicate the shortest path from source.

# Analysis of Acyclic SP


- Topological sort computes a shortest path tree in any edge weighted DAG in time proportional to  $M + N$  (edge weights can be negative!)
  - Each edge is relaxed exactly once (when  $v$  is relaxed), leaving  $\text{distTo}[v] \leq \text{distTo}[u] + w(u, v)$ , so total runtime of acyclic SP is  $M + N + M \approx M + N$
  - Inequality holds until algorithm terminates:
    - $\text{distTo}[v]$  cannot increase because  $\text{distTo}$  values are monotonically decreasing
    - $\text{distTo}[u]$  will not change; no edge pointing to  $u$  will be relaxed after  $u$  is relaxed because of topological order

# Application of Acyclic SP

- Seam carving (Avidan and Shamir): Resize an image for displaying without distortion on a cellphone or web browser
- Enable the user to see the whole image without distortion while scrolling
- Uses DAG shortest path algorithm to find the “shortest path” of pixels through the image (the path that has the lowest energy)
  - The shortest path is almost a column, but not exactly a column

# Content-Aware Resizing

- To find vertical seam, create a DAG of pixels:
  - Vertex = pixel; edge = from pixel to 3 downward neighbors
  - Weight of edge = “energy” (difference in gray levels) of neighboring pixels
  - Seam = shortest path (lowest energy) from top to bottom



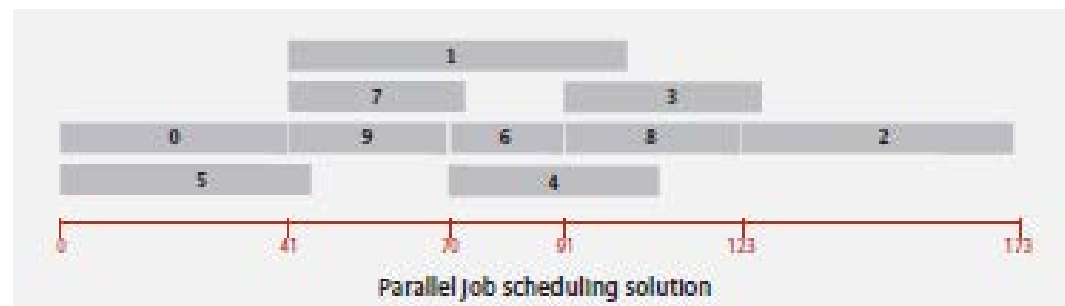
# Acyclic Longest Path Algorithm

- The (acyclic) longest path is called the ***critical path***
- Formulate as an acyclic shortest path problem:
  - Negate all initial weights and run the acyclic shortest path (SP) algorithm as is, or
  - Run acyclic SP, replacing  $\infty$  with  $-\infty$  in the initialize procedure and  $>$  with  $<$  in the relax procedure
- Recall that topological sort algorithm works even with negative weights

# Application of Acyclic LP

- Goal: Given a set of jobs with durations and precedence constraints, find the *minimum* amount of time required for all jobs to complete (i.e., find the bottleneck)
  - Some jobs must be done before others, and some jobs may be performed simultaneously

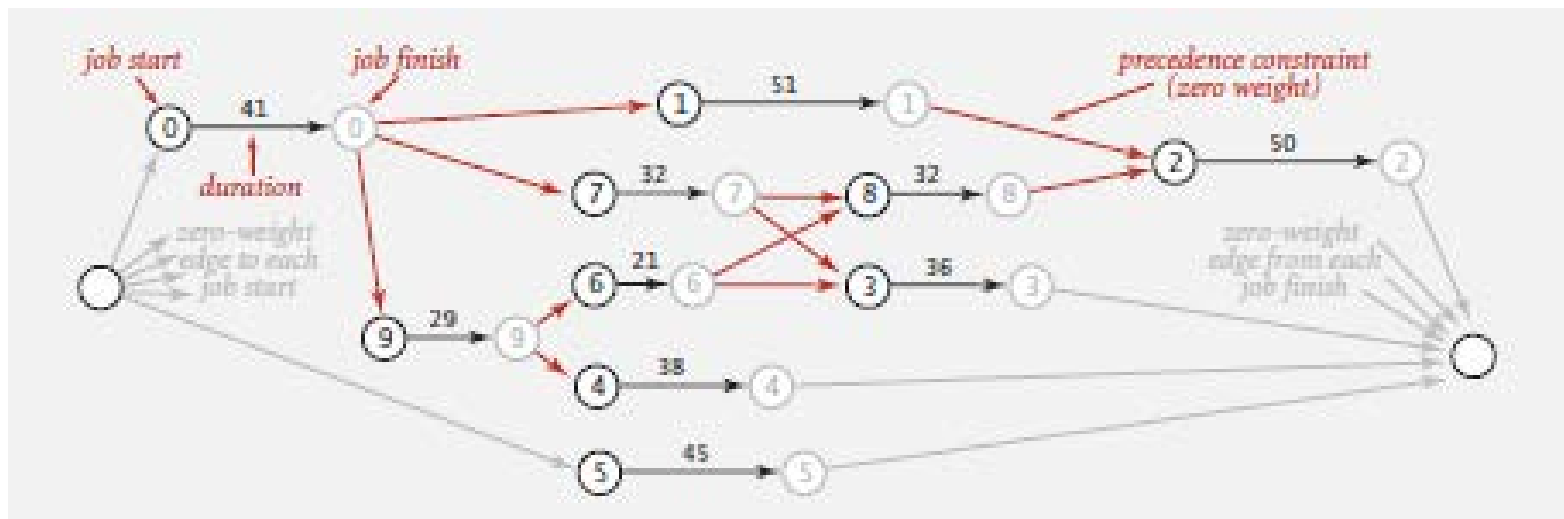
job	duration	must complete before		
0	41.0	1	7	9
1	51.0	2		
2	50.0			
3	36.0			
4	38.0			
5	45.0			
6	21.0	3	8	
7	32.0	3	8	
8	32.0	2		
9	29.0	4	6	



# Application of Acyclic LP

- Create a weighted DAG with source and sink vertices
- Have two vertices (start and finish) for each job
- Have three edges for each job:
  - source to start (0 weight)
  - start to finish (weighted by duration of job)
  - finish to sink (0 weight)
- Have one edge for each precedence constraint (0 weight)

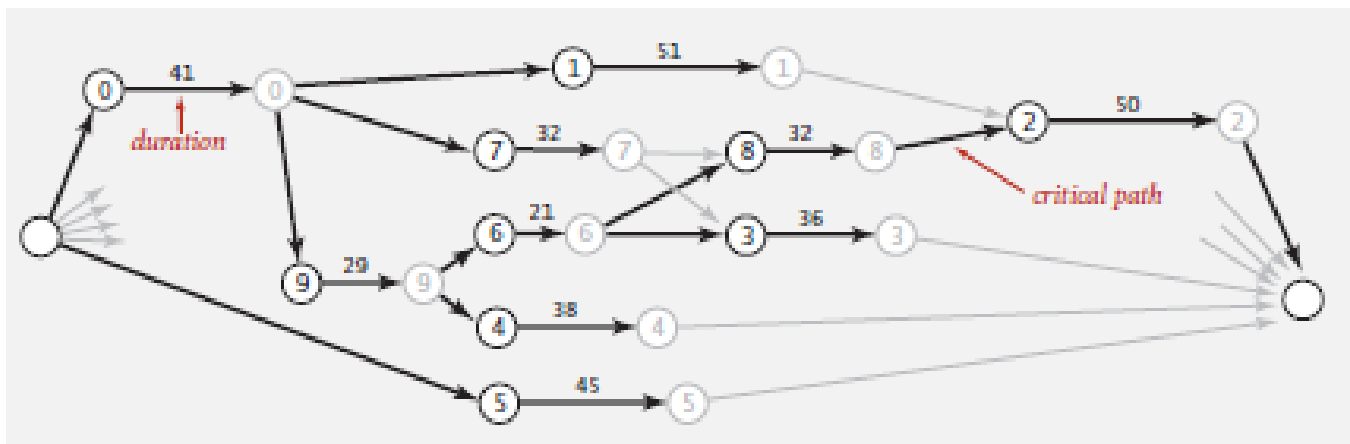
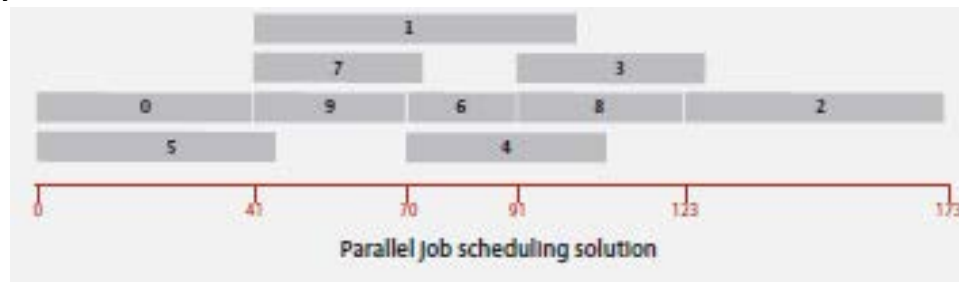
job	duration	must complete before		
0	41.0	1	7	9
1	51.0	2		
2	50.0			
3	36.0			
4	38.0			
5	45.0			
6	21.0	3	8	
7	32.0	3	8	
8	32.0	2		
9	29.0	4	6	





# Application of Acyclic LP

- Now run the “modified” acyclic SP algorithm to get acyclic LP
- The acyclic longest path from the source to the destination is equal to the overall minimum completion time (the bottleneck)



# Difference Constraints

- Goal: optimize a linear function subject to a set of linear inequalities
  - Given an  $M \times N$  matrix  $\mathbf{A}$ , an  $M$ -vector  $\mathbf{b}$ , we wish to find a vector  $\mathbf{x}$  of  $N$  elements that maximizes an objective function, subject to the  $M$  constraints given by  $\mathbf{Ax} \leq \mathbf{b}$
  - This problem can be reduced to finding the shortest paths from a single source

# Difference Constraints

For example, find the 5-element vector  $\mathbf{x}$  that satisfies:


$$\begin{pmatrix} 1 & -1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & -1 \\ 0 & 1 & 0 & 0 & -1 \\ -1 & 0 & 1 & 0 & 0 \\ -1 & 0 & 0 & 1 & 0 \\ 0 & 0 & -1 & 1 & 0 \\ 0 & 0 & -1 & 0 & 1 \\ 0 & 0 & 0 & -1 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{pmatrix} \leq \begin{pmatrix} 0 \\ -1 \\ 1 \\ 5 \\ 4 \\ -1 \\ -3 \\ -3 \end{pmatrix}$$

This problem is equivalent to finding values for the unknowns  $x_1, x_2, x_3, x_4, x_5$  satisfying these 8 difference constraints:

$$\begin{array}{rcl} x_1 - x_2 & \leq & 0 \\ x_1 - x_5 & \leq & -1 \\ x_2 - x_5 & \leq & 1 \\ x_3 - x_1 & \leq & 5 \\ x_4 - x_1 & \leq & 4 \\ x_4 - x_3 & \leq & -1 \\ x_5 - x_3 & \leq & -3 \\ x_5 - x_4 & \leq & -3 \end{array}$$

# Difference Constraints

Create a *constraint graph* with an additional vertex  $v_0$  to guarantee that the graph has a vertex which can reach all other vertices. Include a vertex  $v_i$  for each unknown  $x_i$ . The edge set contains an edge for each difference constraint. Then run the Bellman-Ford algorithm from  $v_0$ .



$$\begin{array}{rcl} x_1 - x_2 & \leq & 0 \\ x_1 - x_5 & \leq & -1 \\ x_2 - x_5 & \leq & 1 \\ x_3 - x_1 & \leq & 5 \\ x_4 - x_1 & \leq & 4 \\ x_4 - x_3 & \leq & -1 \\ x_5 - x_3 & \leq & -3 \\ x_5 - x_4 & \leq & -3 \end{array}$$

One feasible solution to this problem is  $x = (-5, -3, 0, -1, -4)$ .